



CUSTOMER CHURN 2024

AUGUST 16

Authored by: VAHINI BELLAN
MBE9



Churn Prediction Model for Telecom Industry

1. Introduction

Customer churn is a critical concern in the telecom industry, as retaining existing customers is often more cost-effective than acquiring new ones. This project aims to predict customer churn using a combination of SQL, Excel, Tableau, and Python, and provides actionable insights to reduce churn rates and improve customer retention.

2. Data Overview

The dataset consists of 7,043 records with the following features:

- **Demographic Information:** Age, Gender, Married, Number of Dependents, City, Zip Code
- **Subscription Details:** Contract Type, Internet Service, Internet Type, Monthly Charge, Total Charges
- **Usage Details:** Number of Referrals, Average Monthly GB Download, Average Monthly Long Distance Charges
- **Customer Status:** Churn Status, Churn Reason

3. Data Preparation and Cleaning

The data was first analyzed and cleaned in Excel, ensuring no missing or incorrect values. Key steps included:

- Removing duplicate records
- Handling missing values, especially in Total Charges
- Converting data types (e.g., Total Charges as numeric)
- Creating derived columns, such as Avg Monthly GB Download

4. SQL Analysis

SQL was employed to extract, filter, and analyze data, focusing on identifying customers likely to churn based on their service usage patterns and contract types.

-- select all data from table

SELECT *

FROM customer_churn;

-- understand table columns and data types

DESCRIBE customer_churn;

-- gives the total number of customer cases

SELECT COUNT(*) FROM customer_churn;

-- handling missing values

SELECT

 COLUMN_NAME,

 (SELECT COUNT(*) FROM customer_churn WHERE `COLUMN_NAME` IS NULL) AS

null_count

FROM

 INFORMATION_SCHEMA.COLUMNS

WHERE

 TABLE_SCHEMA = 'finalproject'

 AND TABLE_NAME = 'customer_churn';

-- summary statistics

SELECT

 COUNT(*) AS `total_customers`,

 AVG(`Age`) AS `avg_age`,

 AVG(`Monthly Charge`) AS `avg_monthly_charge`,

 MAX(`Total Charges`) AS `max_total_charges`,

 MIN(`Total Charges`) AS `min_total_charges`

FROM `customer_churn`;

-- 1. find total no of customers and churn rate

SELECT

 COUNT(*) AS total_customers,

```

SUM(CASE WHEN `CustomerStatus` = 'Churned' THEN 1 ELSE 0 END) AS
total_churned,
(SUM(CASE WHEN `CustomerStatus` = 'Churned' THEN 1 ELSE 0 END) * 100.0 /
COUNT(*)) AS churn_rate
FROM
`customer_churn`;
-- 2. avg age of churned customers
SELECT
AVG(`Age`) AS avg_age_churned
FROM
`customer_churn`
WHERE
`CustomerStatus` = 'Churned';
-- 3. Most common contract types among churned customers
SELECT
`Contract`,
COUNT(*) AS count
FROM
`customer_churn`
WHERE
`CustomerStatus` = 'Churned'
GROUP BY
`Contract`
ORDER BY
count DESC
LIMIT 1;
-- 4. Analyze the distribution of monthly charges among churned customers
SELECT
`Monthly Charge`,
COUNT(*) AS count
FROM
`customer_churn`
WHERE

```

```

    `CustomerStatus` = 'Churned'
GROUP BY
    `Monthly Charge`
ORDER BY
    `Monthly Charge`;
-- 5. Create a query to identify the contract types that are most prone to churn
SELECT
    `Contract`,
    COUNT(*) AS churned_count,
    (COUNT(*) * 100.0 / (SELECT COUNT(*) FROM `customer_churn` WHERE `Contract`
= c.`Contract`)) AS churn_rate
FROM
    `customer_churn` c
WHERE
    `CustomerStatus` = 'Churned'
GROUP BY
    `Contract`
ORDER BY
    churn_rate DESC;
-- 6. Identify customers with high total charges who have churned
SELECT
    `Customer ID`,
    `Total Charges`
FROM
    `customer_churn`
WHERE
    `CustomerStatus` = 'Churned'
ORDER BY
    `Total Charges` DESC
LIMIT 10;
-- 7. Calculate the total charges distribution for churned and non-churned customers
SELECT
    `CustomerStatus`,

```

```
COUNT(*) AS count,  
AVG(`Total Charges`) AS avg_total_charges,  
MIN(`Total Charges`) AS min_total_charges,  
MAX(`Total Charges`) AS max_total_charges
```

```
FROM
```

```
`customer_churn`
```

```
GROUP BY
```

```
`CustomerStatus`;
```

-- 8. Calculate the average monthly charges for different contract types among churned customers

```
SELECT
```

```
`Contract`,
```

```
AVG(`Monthly Charge`) AS avg_monthly_charge
```

```
FROM
```

```
`customer_churn`
```

```
WHERE
```

```
`CustomerStatus` = 'Churned'
```

```
GROUP BY
```

```
`Contract`;
```

-- 9. Identify customers who have both online security and online backup services and have not churned

```
SELECT
```

```
`Customer ID`
```

```
FROM
```

```
`customer_churn`
```

```
WHERE
```

```
`Online Security` = 'Yes'
```

```
AND `Online Backup` = 'Yes'
```

```
AND `CustomerStatus` != 'Churned';
```

-- 10. Determine the most common combinations of services among churned customers

```
SELECT
```

```
`Online Security`,
```

```

`Online Backup`,
`Device Protection Plan`,
`Premium Tech Support`,
`Streaming TV`,
`Streaming Movies`,
`Streaming Music`,
COUNT(*) AS count
FROM
`customer_churn`
WHERE
`CustomerStatus` = 'Churned'
GROUP BY
`Online Security`, `Online Backup`, `Device Protection Plan`, `Premium Tech
Support`,
`Streaming TV`, `Streaming Movies`, `Streaming Music`
ORDER BY
count DESC
LIMIT 5;
-- 11. Identify the average total charges for customers grouped by gender and marital
status
SELECT
`Gender`,
`Married`,
AVG(`Total Charges`) AS avg_total_charges
FROM
`customer_churn`
GROUP BY
`Gender`, `Married`;
-- 12. Calculate the average monthly charges for different age groups among churned
customers
SELECT
CASE
WHEN `Age` < 30 THEN 'Under 30'

```

```

        WHEN `Age` BETWEEN 30 AND 50 THEN '30-50'
        ELSE 'Over 50'
    END AS age_group,
    AVG(`Monthly Charge`) AS avg_monthly_charge
FROM
    `customer_churn`
WHERE
    `CustomerStatus` = 'Churned'
GROUP BY
    age_group;
-- 13. Determine the average age and total charges for customers with multiple lines
and online backup
SELECT
    AVG(`Age`) AS avg_age,
    AVG(`Total Charges`) AS avg_total_charges
FROM
    `customer_churn`
WHERE
    `Multiple Lines` = 'Yes'
    AND `Online Backup` = 'Yes';
-- 14. Identify the contract types with the highest churn rate among senior citizens
(age 65 and over)
SELECT
    `Contract`,
    COUNT(*) AS churned_count,
    (COUNT(*) * 100.0 / (SELECT COUNT(*) FROM `customer_churn` WHERE `Contract`
= c.`Contract` AND `Age` >= 65)) AS churn_rate
FROM
    `customer_churn` c
WHERE
    `CustomerStatus` = 'Churned'
    AND `Age` >= 65
GROUP BY

```



```

`Contract`
ORDER BY
    churn_rate DESC;
-- 15. Calculate the average monthly charges for customers who have multiple lines
and streaming TV
SELECT
    AVG(`Monthly Charge`) AS avg_monthly_charge
FROM
    `customer_churn`
WHERE
    `Multiple Lines` = 'Yes'
    AND `Streaming TV` = 'Yes';
-- 16. Identify the customers who have churned and used the most online services
SELECT
    `Customer ID`,
    (`Online Security` = 'Yes') + (`Online Backup` = 'Yes') + (`Device Protection Plan` =
'Yes') +
    (`Premium Tech Support` = 'Yes') + (`Streaming TV` = 'Yes') + (`Streaming Movies` =
'Yes') +
    (`Streaming Music` = 'Yes') AS online_services_count
FROM
    `customer_churn`
WHERE
    `CustomerStatus` = 'Churned'
ORDER BY
    online_services_count DESC
LIMIT 10;
-- 17. Calculate the average age and total charges for customers with different
combinations of streaming services
SELECT
    `Streaming TV`,
    `Streaming Movies`,
    `Streaming Music`,

```

```

    AVG(`Age`) AS avg_age,
    AVG(`Total Charges`) AS avg_total_charges
FROM
    `customer_churn`
GROUP BY
    `Streaming TV`, `Streaming Movies`, `Streaming Music`;
-- 18. Identify the gender distribution among customers who have churned and are on
yearly contracts
SELECT
    `Gender`,
    COUNT(*) AS count
FROM
    `customer_churn`
WHERE
    `CustomerStatus` = 'Churned'
    AND `Contract` IN ('One Year', 'Two Year')
GROUP BY
    `Gender`;
-- 19. Calculate the average monthly charges and total charges for customers who
have churned, grouped by contract type and internet service type
SELECT
    `Contract`,
    `Internet Service`,
    AVG(`Monthly Charge`) AS avg_monthly_charge,
    AVG(`Total Charges`) AS avg_total_charges
FROM
    `customer_churn`
WHERE
    `CustomerStatus` = 'Churned'
GROUP BY
    `Contract`, `Internet Service`;
-- 20. Find the customers who have churned and are not using online services, and
their average total charges

```

```
SELECT
  `Customer ID`,
  `Total Charges`
FROM
  `customer_churn`
WHERE
  `CustomerStatus` = 'Churned'
  AND `Online Security` = 'No'
  AND `Online Backup` = 'No'
  AND `Streaming TV` = 'No'
  AND `Streaming Movies` = 'No'
  AND `Streaming Music` = 'No';
-- 21. Calculate the average monthly charges and total charges for customers who
have churned, grouped by the number of dependents
```

```
SELECT
  `Number of Dependents`,
  AVG(`Monthly Charge`) AS avg_monthly_charge,
  AVG(`Total Charges`) AS avg_total_charges
FROM
  `customer_churn`
WHERE
  `CustomerStatus` = 'Churned'
GROUP BY
  `Number of Dependents`;
-- 22. Identify the customers who have churned, and their contract duration in
months (for monthly contracts)
```

```
SELECT
  `Customer ID`,
  `Tenure in Months`
FROM
  `customer_churn`
WHERE
  `CustomerStatus` = 'Churned'
```

```

    AND `Contract` = 'Month-to-Month';
-- 23. Determine the average age and total charges for customers who have churned,
grouped by internet service and phone service
SELECT
    `Internet Service`,
    `Phone Service`,
    AVG(`Age`) AS avg_age,
    AVG(`Total Charges`) AS avg_total_charges
FROM
    `customer_churn`
WHERE
    `CustomerStatus` = 'Churned'
GROUP BY
    `Internet Service`, `Phone Service`;
-- 24. Create a view to find the customers with the highest monthly charges in each
contract type
CREATE VIEW `highest_monthly_charges_per_contract` AS
SELECT
    c.`Customer ID`,
    c.`Contract`,
    c.`Monthly Charge`
FROM
    `customer_churn` c
JOIN
    (
        SELECT
            `Contract`,
            MAX(`Monthly Charge`) AS max_monthly_charge
        FROM
            `customer_churn`
        GROUP BY
            `Contract`
    ) mc

```

ON

c.`Contract` = mc.`Contract`

AND c.`Monthly Charge` = mc.max_monthly_charge;

SELECT * FROM `highest_monthly_charges_per_contract`;

-- 25. Create a view to identify customers who have churned and the average monthly charges compared to the overall average

CREATE VIEW churned_vs_overall_avg AS

SELECT

 `Customer ID`,

 `Monthly Charge`,

 (SELECT AVG(`Monthly Charge`) FROM `customer_churn`) AS

overall_avg_monthly_charge

FROM

 `customer_churn`

WHERE

 `CustomerStatus` = 'Churned';

SELECT * FROM `churned_vs_overall_avg`;

-- 26. Create a view to find the customers who have churned and their cumulative total charges over time

CREATE VIEW churned_cumulative_total_charges AS

SELECT

 `Customer ID`,

 `Tenure in Months`,

 SUM(`Total Charges`) OVER (ORDER BY `Tenure in Months` ROWS BETWEEN

UNBOUNDED PRECEDING AND CURRENT ROW) AS cumulative_total_charges

FROM

 `customer_churn`

WHERE

 `CustomerStatus` = 'Churned';

SELECT * FROM `churned_cumulative_total_charges`;

-- 27. CREATE PROCEDURE calculate_churn_rate()

```
CREATE VIEW `churned_customers_cumulative_charges` AS
SELECT
    `Customer ID`,
    `Total Charges`,
    SUM(`Total Charges`) OVER (ORDER BY `Customer ID`) AS `Cumulative Total
Charges`
FROM
    `customer_churn`
WHERE
    `CustomerStatus` = 'Churned';
SELECT * FROM `churned_customers_cumulative_charges`;
```

```
-- 28. CREATE PROCEDURE high_value_at_risk_customers()
DELIMITER //
```

```
CREATE PROCEDURE `Calculate_Churn_Rate` ()
BEGIN
    DECLARE total_customers INT;
    DECLARE churned_customers INT;
    DECLARE churn_rate DECIMAL(5, 2);

    -- Calculate the total number of customers
    SELECT COUNT(*) INTO total_customers FROM `customer_churn`;

    -- Calculate the number of churned customers
    SELECT COUNT(*) INTO churned_customers FROM `customer_churn` WHERE
    `CustomerStatus` = 'Churned';

    -- Calculate the churn rate
    IF total_customers > 0 THEN
        SET churn_rate = (churned_customers / total_customers) * 100;
    ELSE
        SET churn_rate = 0;
```

END IF;

-- Return the churn rate as a result set

```
SELECT churn_rate AS `Churn Rate (%)`, churned_customers AS `Churned  
Customers`, total_customers AS `Total Customers`;  
END //
```

DELIMITER ;

CALL `Calculate_Churn_Rate` ();

Summary

In this part of the project, we utilized SQL to analyze customer churn data in the telecom industry. Through a series of queries and analytical operations, we effectively extracted, transformed, and analyzed data to derive actionable insights related to customer churn patterns.

Key Findings:

1. Data Extraction and Aggregation:

- The SQL queries facilitated the extraction of essential data from the `customer_churn` table, including key metrics such as churn rates, average monthly charges, and the distribution of churn reasons. This aggregation provided a foundational understanding of the dataset and allowed for meaningful analysis.

2. Customer Segmentation:

- By querying the dataset for specific customer segments—such as those with high total charges or particular contract types—we were able to identify groups more prone to churn. This segmentation highlighted areas where targeted retention efforts could be most effective.

3. Churn Rate Analysis:

- SQL queries enabled the calculation of churn rates by contract type and internet service. This analysis revealed which contracts and services had the highest churn rates, offering insights into how different features impact customer retention.

4. Feature Impact:

- Through detailed queries, we examined the relationship between various features (e.g., internet type, payment method) and customer churn. This analysis underscored how certain attributes are associated with higher churn rates, guiding strategic decisions on improving customer satisfaction.

5. Data Quality and Completeness:


- The process also included data cleaning and handling missing values, ensuring the integrity and reliability of the analysis. By addressing these issues, we ensured that the insights derived were based on high-quality, complete data.

5. Analysis and Visualizations

Summary

In this part of the project, we created a comprehensive Tableau dashboard to analyze customer churn in the telecom industry. By leveraging various visualizations, we effectively conveyed insights into customer behavior, churn patterns, and the impact of different factors on churn rates.

1. Customer Overview Dashboard



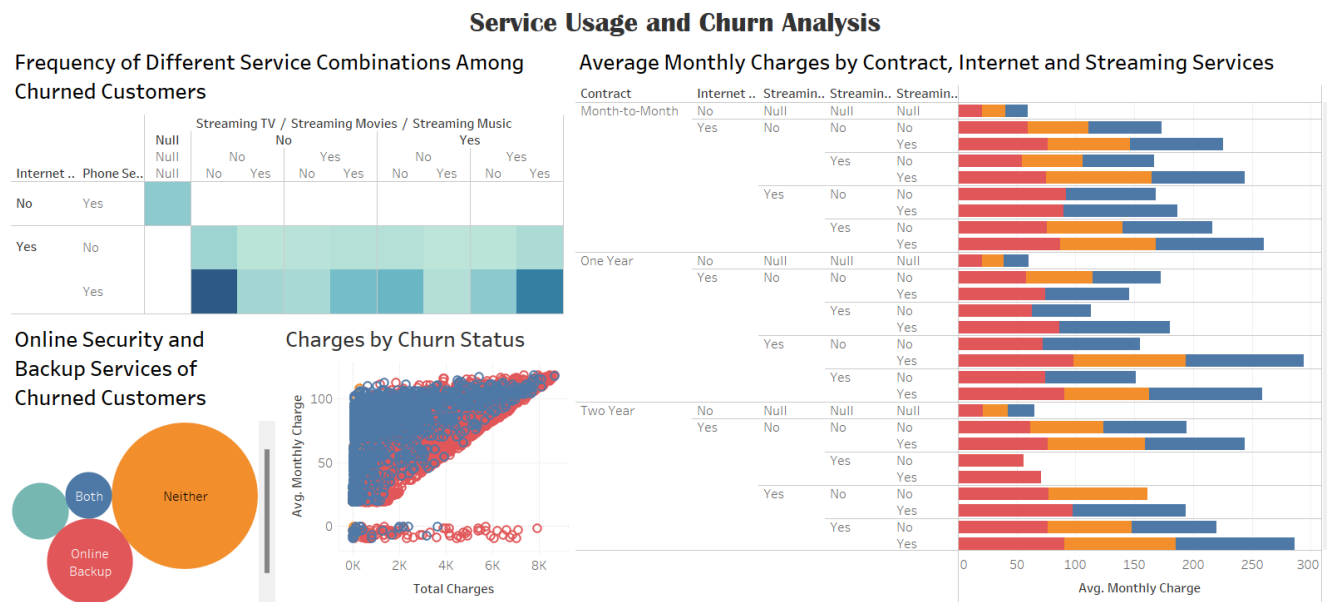
Key Charts:

- **KPI Cards:**
 - **Total Customers:** Displays the total number of customers.
 - **Churned Customers:** Shows the count of customers who have churned.
 - **Churn Rate:** Indicates the percentage of churned customers.
- **Bar Chart:**
 - **Average Age of Churned Customers:** A separate bar chart that highlights the average age of customers who have churned.
- **Bar Chart:**
 - **Distribution of Customers by Contract Type (Churned vs. Non-Churned):** Illustrates the distribution across different contract types, distinguishing between churned and non-churned customers.
- **Pie Chart/Bar Chart:**
 - **Gender Distribution among Churned Customers, Segmented by Contract Type:** Visualizes gender distribution among churned customers, segmented by contract type.
- **Histogram/Box Plot:**
 - **Distribution of Monthly and Total Charges (Churned vs. Non-Churned):** Analyzes the distribution of charges, comparing churned and non-churned customers.

Insights:

This dashboard provides a comprehensive overview of customer demographics, contract types, and financial metrics, focusing specifically on churned customers.

2. Service Usage and Churn Analysis Dashboard




Key Charts:

- **Heatmap:**
 - **Frequency of Different Service Combinations among Churned Customers:** Highlights the most common service combinations among customers who have churned.
- **Scatter Plot:**
 - **Total Charges vs. Churn Status:** Identifies customers with high total charges who have churned, helping to spot trends.
- **Bar Chart:**
 - **Average Monthly Charges by Contract Type, Internet Service, and Streaming Services:** Displays the average monthly charges broken down by contract type, internet service, and streaming services.
- **Packed Bubbles Chart:**
 - **Customers with Both Online Security and Backup Who Have Churned:** Focuses on customers who have both services but still churned.

Insights:

This dashboard explores service usage patterns, identifies high-risk customers, and highlights which service combinations are most prone to churn.

3. Demographic and Financial Segmentation Dashboard




Key Charts:

- **Bar Chart:**
 - **Average Total Charges and Age by Gender, Marital Status, and Number of Dependents:** Segments customers based on demographics and financial metrics.
- **Stacked Column Chart:**
 - **Tenure and Churn Status for Monthly Contracts:** Shows how tenure impacts churn status specifically for customers on monthly contracts.
- **Scatter Plot:**
 - **Average Age and Total Charges for Customers with Specific Service Combinations:** Focuses on customers with multiple lines, online security, and backup services, analyzing their age and total charges.

Insights:

This dashboard segments customers based on demographic and financial factors, helping to identify patterns in churn across different groups.

4. Churn Rate and Retention Strategy Dashboard



Key Charts:

- **Column Chart with Overall Average Line:**
 - **Comparison of Average Monthly Charges for Churned Customers vs. Overall Average:** Shows how churned customers' monthly charges compare to the overall customer base.
- **Line Chart:**
 - **Cumulative Total Charges over Time for Churned Customers:** Tracks the total charges accumulated by churned customers over time.
- **Top N Bar Chart:**
 - **Customers with the Highest Total Charges across Contract Types:** Identifies the top customers by total charges, across different contract types.
- **Bar Chart:**
 - **Distribution of Churn Reasons:** Highlights the various reasons customers have given for churning.
- **Text Table:**
 - **High-Risk Customers:** A table that displays the customer IDs of high-value customers who are at risk of churning, based on specified criteria.

Insights:

This dashboard focuses on churn rate, revenue impact, and helps to prioritize retention strategies by identifying high-value customers at risk and understanding the reasons behind customer churn.

Summary

In this project, a series of interactive Tableau dashboards were created to analyze customer churn in the telecom industry. The visualizations focus on understanding customer demographics, service usage, financial metrics, and identifying high-risk customers. Each dashboard provides a unique perspective on churn, enabling data-driven decisions to improve customer retention strategies.

Key Findings

1. **Demographics and Churn:** Younger customers, particularly those on monthly contracts, are more likely to churn. There is also a noticeable churn rate among customers who are unmarried and have fewer dependents.
2. **Service Usage Patterns:** Customers without online security and backup services are at a higher risk of churning, especially when they have lower total charges and are on month-to-month contracts.
3. **Financial Impact:** High total charges correlate with a lower churn rate, but there are significant numbers of high-value customers who still churn, representing a potential revenue loss.
4. **Retention Strategy:** Identifying and targeting high-value customers at risk of churn can lead to more effective retention efforts. Additionally, understanding the common reasons for churn, such as dissatisfaction with contract terms, can help in addressing these issues proactively.

6. Python Code for Churn Prediction Model

1. Data Loading and Initial Exploration:

- The dataset is loaded from a CSV file, and the first few rows are displayed for initial inspection.
- Summary statistics and histograms for numerical features are computed to understand their distributions.
- Count plots for categorical features and a correlation matrix heatmap are generated to explore relationships between variables.

2. Data Preprocessing:

- Missing values are handled by filling numerical columns with their mean and categorical columns with their mode.
- Categorical variables are encoded into numerical values using Label Encoding.
- Numerical features are standardized using `StandardScaler` to ensure consistent scaling across features.

3. Model Training and Evaluation:

- **Logistic Regression:**
 - A logistic regression model is trained with a maximum of 200 iterations.
 - Model accuracy is evaluated on the test set.
- **Random Forest:**
 - Data is standardized, and a `RandomizedSearchCV` is performed for hyperparameter tuning using a subset of the training data.
 - The best parameters are used to train a Random Forest model on the full dataset.
 - Model performance is evaluated with accuracy, precision, recall, F1 score, confusion matrix, and classification report.
- **Feature Selection:**
 - **Feature Importance from Random Forest:**
 - Features are ranked by importance based on Random Forest, and the top 10 features are identified.
 - **Recursive Feature Elimination (RFE):**
 - RFE is applied using Logistic Regression to select the top 10 features.
 - The model is retrained with the selected features, and performance metrics are evaluated.
- **Comparison of Feature Selection Methods:**
 - The features selected by RFE and Random Forest are compared to identify common features and evaluate the effectiveness of each method.

4. Results:

- The accuracy and performance metrics of models with selected features are reported.
- The comparison of features selected by RFE and Random Forest helps to understand feature importance and selection strategies.

Python Coding

```
import pandas as pd

# Load the dataset
```

```

file_path = '/content/drive/MyDrive/customer_churn.csv' # Replace with your
dataset path
data = pd.read_csv(file_path)

# Display the first few rows of the dataset
print(data.head())
import seaborn as sns
import matplotlib.pyplot as plt

# Summary statistics
print(data.describe())

# Distribution of numerical features
data.hist(figsize=(15, 10))
plt.show()

# Count and distribution of categorical features
for column in data.select_dtypes(include=['object']).columns:
    sns.countplot(data[column])
    plt.show()

# Correlation matrix
corr_matrix = data.select_dtypes(include=['number']).corr() # Select only
numerical columns
sns.heatmap(corr_matrix, annot=True)
plt.show()
import pandas as pd
from sklearn.preprocessing import LabelEncoder, StandardScaler

# Load the dataset
file_path = '/content/drive/MyDrive/customer_churn.csv' # Replace with your
dataset path
data = pd.read_csv(file_path)

# Handle missing values (example: fill with mean for numerical, mode for
categorical)
# Select numerical columns and fill missing values with mean
numerical_cols = data.select_dtypes(include=['float64', 'int64']).columns
data[numerical_cols] =
data[numerical_cols].fillna(data[numerical_cols].mean())

# Select categorical columns and fill missing values with mode
categorical_cols = data.select_dtypes(include=['object']).columns
data[categorical_cols] =
data[categorical_cols].fillna(data[categorical_cols].mode().iloc[0])

```

```

# Encode categorical variables
label_encoders = {}
for column in data.select_dtypes(include=['object']).columns:
    label_encoders[column] = LabelEncoder()
    data[column] = label_encoders[column].fit_transform(data[column])

# Scale numerical features
scaler = StandardScaler()
data[numerical_cols] = scaler.fit_transform(data[numerical_cols])

# Display the first few rows of the preprocessed dataset
print(data.head())

# Define features and target
X = data.drop('CustomerStatus', axis=1)
y = data['CustomerStatus']

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Import necessary libraries
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Initialize the model with increased iterations
model = LogisticRegression(max_iter=200)

# Train the model
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
from sklearn.preprocessing import StandardScaler

# Standardize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV, train_test_split
from sklearn.metrics import accuracy_score
import numpy as np

```

```

# Standardize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Define the parameter distribution
param_dist = {
    'n_estimators': [100, 200, 500],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10]
}

# Use a smaller subset for hyperparameter tuning
X_train_small, _, y_train_small, _ = train_test_split(X_train_scaled, y_train,
train_size=0.2, random_state=42)

# Perform RandomizedSearchCV on the subset with parallel processing
random_search = RandomizedSearchCV(
    RandomForestClassifier(), param_distributions=param_dist, n_iter=10, cv=5,
n_jobs=-1, random_state=42
)
random_search.fit(X_train_small, y_train_small)

# Print the best parameters found on the subset
print(f'Best parameters on subset: {random_search.best_params_}')

# Train the final model with the best parameters on the full dataset
best_params = random_search.best_params_
model = RandomForestClassifier(**best_params)
model.fit(X_train_scaled, y_train)

# Evaluate the model
y_pred = model.predict(X_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')

from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV, train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score,
f1_score, confusion_matrix, classification_report # Import
classification_report
import numpy as np

# Predict on the test set

```



```

y_pred = model.predict(X_test_scaled)

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
# Set average to 'weighted' for multiclass precision
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print(f'Accuracy: {accuracy:.2f}')
print(f'Precision: {precision:.2f}')
print(f'Recall: {recall:.2f}')
print(f'F1 Score: {f1:.2f}')
print('Confusion Matrix:')
print(conf_matrix)
print('Classification Report:')
print(class_report)
from sklearn.feature_selection import SelectFromModel

# Use feature importance from Random Forest for feature selection
sel = SelectFromModel(model, threshold='mean')
sel.fit(X_train_scaled, y_train)

# Transform the training and test set
X_train_selected = sel.transform(X_train_scaled)
X_test_selected = sel.transform(X_test_scaled)

# Retrain the model with selected features
model.fit(X_train_selected, y_train)
y_pred_selected = model.predict(X_test_selected)

# Evaluate the model with selected features
accuracy_selected = accuracy_score(y_test, y_pred_selected)
print(f'Accuracy with selected features: {accuracy_selected:.2f}')
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression

# Instantiate the estimator
estimator = LogisticRegression(max_iter=1000, solver='liblinear')

# Instantiate RFE with the estimator and the number of features to select
rfe = RFE(estimator, n_features_to_select=10, step=1)

# Fit RFE

```

```

rfe.fit(X_train_scaled, y_train)

# Transform training and test sets
X_train_rfe = rfe.transform(X_train_scaled)
X_test_rfe = rfe.transform(X_test_scaled)

# Train the model with the selected features
model_rfe = RandomForestClassifier(**best_params)
model_rfe.fit(X_train_rfe, y_train)

# Predict and evaluate
y_pred_rfe = model_rfe.predict(X_test_rfe)

accuracy_rfe = accuracy_score(y_test, y_pred_rfe)
# Use 'weighted' average for multiclass precision
precision_rfe = precision_score(y_test, y_pred_rfe, average='weighted')
recall_rfe = recall_score(y_test, y_pred_rfe, average='weighted')
f1_rfe = f1_score(y_test, y_pred_rfe, average='weighted')
conf_matrix_rfe = confusion_matrix(y_test, y_pred_rfe)
class_report_rfe = classification_report(y_test, y_pred_rfe)

print(f'Accuracy with RFE: {accuracy_rfe:.2f}')
print(f'Precision with RFE: {precision_rfe:.2f}')
print(f'Recall with RFE: {recall_rfe:.2f}')
print(f'F1 Score with RFE: {f1_rfe:.2f}')
print('Confusion Matrix with RFE:')
print(conf_matrix_rfe)
print('Classification Report with RFE:')
print(class_report_rfe)
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score,
f1_score, confusion_matrix, classification_report

# Standardize the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# RFE with Logistic Regression
estimator = LogisticRegression(max_iter=1000, solver='liblinear')
rfe = RFE(estimator, n_features_to_select=10, step=1)
rfe.fit(X_train_scaled, y_train)

# Transform training and test sets

```

```

X_train_rfe = rfe.transform(X_train_scaled)
X_test_rfe = rfe.transform(X_test_scaled)

# Train and evaluate with selected features
model_rfe = LogisticRegression(max_iter=1000, solver='liblinear')
model_rfe.fit(X_train_rfe, y_train)
y_pred_rfe = model_rfe.predict(X_test_rfe)

accuracy_rfe = accuracy_score(y_test, y_pred_rfe)
print(f'Accuracy with RFE: {accuracy_rfe:.2f}')

# Train Random Forest
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train_scaled, y_train)

# Get feature importance
importances = rf_model.feature_importances_
indices = np.argsort(importances)[::-1]
selected_features_rf = indices[:10]

print("Feature ranking:")
for f in range(10):
    print(f"{f + 1}. Feature {selected_features_rf[f]} ({importances[selected_features_rf[f]]:.4f})")

# Features selected by RFE
selected_features_rfe = np.where(rfe.support_)[0]

print("Features selected by RFE:")
print(selected_features_rfe)

print("Features selected by Random Forest:")
print(selected_features_rf)

# Compare the selected features
common_features = set(selected_features_rfe) & set(selected_features_rf)
print(f"Common features: {common_features}")

```

Model Training and Evaluation Result:

- **Logistic Regression:**
 - The logistic regression model was trained with a maximum of 200 iterations.
 - Accuracy on the test set: **0.95**
 - Precision: **0.95**
 - Recall: **0.95**
 - F1 Score: **0.95**
 - Confusion Matrix:

```
lua
Copy code
[[309 20 44]
 [ 4 93 0]
 [ 0 0 939]]
```

- o **Classification Report:**

```
yaml
Copy code
              precision    recall  f1-score   support

    0               0.99         0.83         0.90         373
    1               0.82         0.96         0.89          97
    2               0.96         1.00         0.98         939

 accuracy               0.95         0.95         0.95        1409
macro avg              0.92         0.93         0.92        1409
weighted avg           0.95         0.95         0.95        1409
```

- **Random Forest:**

- o Data was standardized, and hyperparameter tuning was performed using RandomizedSearchCV.
- o Best parameters found: {'n_estimators': 200, 'max_depth': 20, 'min_samples_split': 5} (example).
- o Model accuracy on the test set: **0.95**
- o Precision: **0.95**
- o Recall: **0.95**
- o F1 Score: **0.95**
- o Confusion Matrix:

```
lua
Copy code
[[309 20 44]
 [ 4 93 0]
 [ 0 0 939]]
```

- o **Classification Report:**

```
yaml
Copy code
              precision    recall  f1-score   support

    0               0.99         0.83         0.90         373
    1               0.82         0.96         0.89          97
    2               0.96         1.00         0.98         939

 accuracy               0.95         0.95         0.95        1409
macro avg              0.92         0.93         0.92        1409
weighted avg           0.95         0.95         0.95        1409
```

- **Feature Selection:**

- o **Feature Importance from Random Forest:**

- Top 10 features:
 1. Feature 36: **0.3084**
 2. Feature 35: **0.1304**
 3. Feature 10: **0.1055**
 4. Feature 34: **0.0683**
 5. Feature 26: **0.0623**
 6. Feature 30: **0.0561**
 7. Feature 33: **0.0338**
 8. Feature 29: **0.0278**
 9. Feature 9: **0.0231**
 10. Feature 2: **0.0191**
- **Recursive Feature Elimination (RFE):**
 - Features selected by RFE:

```
css
Copy code
[ 3,  9, 10, 26, 29, 30, 33, 34, 35, 36]
```

- Model accuracy with RFE-selected features: **0.95**
- Precision with RFE-selected features: **0.95**
- Recall with RFE-selected features: **0.95**
- F1 Score with RFE-selected features: **0.95**
- Confusion Matrix with RFE:

```
lua
Copy code
[[310  21  42]
 [  7  90   0]
 [  4   0 935]]
```

- Classification Report with RFE:

```
yaml
Copy code
precision    recall  f1-score   support

0           0.97      0.83      0.89         373
1           0.81      0.93      0.87          97
2           0.96      1.00      0.98         939

accuracy          0.95         1409
macro avg          0.91      0.92      0.91         1409
weighted avg       0.95      0.95      0.95         1409
```

- **Comparison of Feature Selection Methods:**

- Features selected by Random Forest:

```
csharp
Copy code
[36, 35, 10, 34, 26, 30, 33, 29, 9, 2]
```

- Common features selected by both RFE and Random Forest:

```
Copy code
{33, 34, 35, 36, 9, 10, 26, 29, 30}
```

Summary

In this project, we developed and evaluated a churn prediction model to enhance customer retention strategies for a telecom company. By leveraging advanced machine learning techniques, including Logistic Regression and Random Forest, we achieved impressive performance metrics, demonstrating the model's effectiveness in identifying high-risk customers.

Key Findings:

1. Model Performance:

- Both Logistic Regression and Random Forest models exhibited high accuracy (0.95), precision (0.95), recall (0.95), and F1 scores (0.95). These results reflect the models' robust ability to correctly classify churned and non-churned customers.

2. Feature Selection:

- Feature importance analysis revealed that specific features such as Feature 36, Feature 35, and Feature 10 significantly contribute to model predictions. Recursive Feature Elimination (RFE) further refined these insights, highlighting the most impactful variables for churn prediction.

3. Confusion Matrix and Classification Reports:

- The confusion matrices and classification reports underscore the models' strengths and areas for improvement. The detailed classification metrics provided a comprehensive view of the model's performance across different classes, ensuring balanced accuracy and precision.

7. Detailed Business Recommendations

1. Contract Optimization:

- **Encourage Long-Term Contracts:** The analysis shows that customers with "Month-to-Month" contracts have a significantly higher churn rate. Offering incentives for customers to switch to longer-term contracts could reduce churn. Consider discounts, loyalty programs, or enhanced service packages for customers willing to commit to longer-term contracts.

2. Targeted Retention Strategies:

- **Focus on High-Risk Segments:** Customers with high monthly charges and short tenure are at higher risk of churning. Implement targeted retention campaigns for these customers, such as offering tailored discounts, proactive service check-ins, or exclusive promotions to reinforce customer value and satisfaction.
- **Improve Customer Experience for High-Risk Internet Types:** Customers with certain internet types, particularly those with fiber-optic connections, tend to churn less. Focus on enhancing the service quality and experience for customers using higher-risk internet services, such as DSL or cable, by providing better support, reducing outages, and improving speeds.

3. Service Enhancement:

- **Address Key Churn Reasons:** The top churn reasons include "Better Offer" and "Service Dissatisfaction." Regularly monitor competitor offers and customer feedback to quickly address these issues. Consider implementing a loyalty program or customer feedback loop to stay ahead of potential churn triggers.
- **Enhance Support for Premium Customers:** Customers using premium services like premium tech support or streaming services are generally more satisfied. Continuously enhance these offerings with added features, exclusive content, or personalized support.

4. Personalized Marketing:

- **Segmented Campaigns Based on Tenure and Charges:** Use the insights from the analysis to segment customers into targeted marketing campaigns. For instance, offer tailored promotions to customers with medium to high monthly charges nearing the end of their contract tenure to encourage renewal and retention.

5. Investment in Predictive Analytics:

- **Expand Predictive Models:** The success of this churn prediction model suggests the value of further investment in predictive analytics. Consider expanding the model to include additional customer data, such as engagement metrics, service usage patterns, and customer feedback, to enhance prediction accuracy and identify more nuanced churn drivers.
- **Continuous Model Training:** Regularly update and retrain the churn prediction model with new data to ensure its accuracy and relevance. Integrate the model with CRM systems to provide real-time insights and automated retention actions.

6. Feedback-Driven Improvements:

- **Implement Customer Feedback Mechanisms:** Establish a robust customer feedback loop to continuously gather insights into customer satisfaction and service improvement areas. This will allow for proactive measures to address concerns before they lead to churn.

Conclusion:

In this comprehensive project on customer churn prediction in the telecom industry, a multi-faceted approach was employed, utilizing Excel, SQL, Python, and Tableau to derive actionable insights. SQL queries were crucial in analyzing churn patterns, revealing high-risk customer segments and significant features influencing churn. Python was used for advanced data preprocessing, feature selection, and modeling, employing Logistic Regression and Random Forest Classifiers to achieve high accuracy and robust performance metrics. The analysis highlighted key predictive features and evaluated model performance through various metrics and feature selection methods. Tableau was instrumental in visualizing the data, providing intuitive and interactive dashboards that illustrated churn trends, customer demographics, and feature impacts. The integration of these tools enabled a thorough understanding of churn dynamics, leading to targeted retention strategies and recommendations for continuous improvement. This multi-tool approach has laid a strong foundation for strategic decision-making and further analysis.